

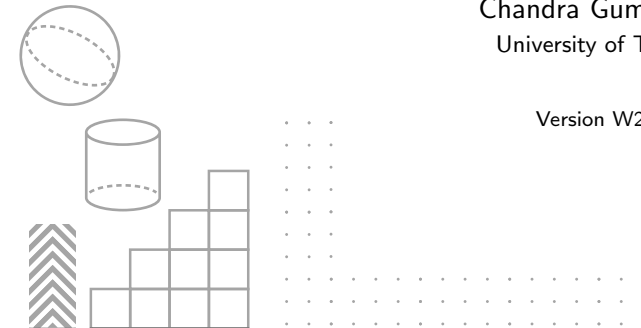
Chapter 6

Search: Adversarial Search

Introduction to Artificial Intelligence

Chandra Gummaluru
University of Toronto

Version W22.1




- The following is based on material developed by many individuals, including (but not limited to):
 - Sheila McIlraith
 - Bahar Aameri
 - Fahiem Bacchus
 - Sonya Allin

- Recall the formal definition of a search problem.
- **Definition:** Search Problem
 - Let \mathcal{S} be a set of **states** that we want to search through.
 - From any given state, $s \in \mathcal{S}$, there exist a set of **actions**, $A(s)$.
 - When an action, $a \in A(s)$, is applied to s , the result is a new state, denoted $a(s)$.
 - A sequence of actions, $\langle a_1, \dots, a_n \rangle$ defines a **path** between two states.
 - The **length** of the path is the number of actions that make it up.
 - Each action, a , may have an associated **cost**, $c(a) > 0$. In this case, the cost of the path is the cumulative cost of its actions.
 - Given some initial state, s_0 , we seek a path (often the shortest/cheapest one) to some state in a subset, $\mathcal{G} \subseteq \mathcal{S}$, called the **goal space**.

Formalizing a Search Problem: Search Trees

- We saw that a graph could be used to represent the structure of a search problem. However, the searching itself takes place on a tree of possible paths.

Example: Search Tree for 4-Queens Puzzle



Formalizing a Search Problem: Assumptions

- So far, we have assumed our agent has complete control over the states.
- In particular, we assumed:
 - ① Our agent is the only actor.
 - ② The actions have deterministic consequences.
- However, these are very strong assumptions.
- We need to modify our algorithms to work even when the aforementioned assumptions do not hold.

- If the first assumption is violated, our agent is said to be playing a game.
- A **game** is a situation in which:
 - There are multiple agents capable of making changes to the environment.
 - Each agent has their own goals and tries to alter the environment to its own benefit.
- Using search to make decisions in a game is difficult because the optimal action from any given state depends on the future actions of other agents.


- In general, a game can be:
 - ① for any number of players, say N .
 - ② **deterministic** (i.e., do not involve randomness), or **stochastic** (i.e., not deterministic).
 - ③ **finite** (i.e., the state-space and action sets are finite), or **infinite** (i.e., not finite).
 - ④ **discrete** (i.e., state-space is discrete), or **continuous** (i.e., not discrete).
 - ⑤ of **perfect information** (i.e., all states are either fully observable by all players), or **partial information** (i.e., not of perfect information).
 - ⑥ **zero-sum** (i.e., the cumulative utility over all players is independent of the state), or **non-zero-sum** (i.e., not zero-sum).

- For simplicity, we will focus on 2-player, deterministic, finite, discrete, zero-sum games of perfect information.
- **Definition:** Components of a Game
 - Let M denote our agent and m denote the adversary.
 - Let \mathcal{S} be the set of **states** that the game could be in.
 - A **circumstance**, k , is a state-player pair, i.e., game's state, and the turn-taker.
 - For any given circumstance, k , there is a set of **actions**, $A(k)$, that must be made by the turn-taker.
 - When an action, $a \in A(k)$, is applied to k , the result is a new circumstance, $a(k)$.
 - We assume that the turn-taking player switches after each action.
 - Let $\mathcal{T} \subseteq \mathcal{S}$ denote the set of **terminal states**, i.e., those in which the game ends.
 - Each terminal state, $t \in \mathcal{T}$, provides a utility, $u(t)$ to M and $-u(t)$ to m .
 - Given a circumstance, k , in which M is the turn-taking player, we seek an action, $a \in A(k)$, that (eventually) leads to a terminal state, $t \in \mathcal{T}$, which maximizes $u(t)$.

Formalizing a Game: Game Trees

- Like any search problem, we can create a tree of possible paths. However, in a game tree, each level corresponds to different player, namely, the turn-taker.

Example: Game Tree for Tic-Tac-Toe Puzzle



- Given any circumstance, $k = (s, M)$, in which M is the turn-taker, they should chose the action, $a_M(s)$, that eventually leads to a terminal state with the maximum utility.
- If s is “almost-terminal”, i.e., one action away from being terminal, the choice is straight-forward:


$$a_M(s) = \operatorname{argmax}_{a \in A(s)} \{u(a(s))\}$$

- If s is not “almost-terminal”, $u(a(s))$, depends on m ’s actions.
- If M knew that m will play $a_m(s)$, in any circumstance, $k = (s, m)$, in which it is the turn-taker, then:

$$U(s, t) = \begin{cases} u(s), s \in \mathcal{T} \\ U(a_m(s), M), t = m \\ \max \{U(s', m), s' \in S(s)\}, t = M \end{cases}$$

Game Strategies: Example

- We will analyse different strategies on the game tree shown below, where \bullet and \circ respectively, denote circumstances in which the turn-taker is M and m :



- In the **min-max** strategy, we assume m always plays its best-response, i.e.,


$$a_m(s) = \operatorname{argmin}_{a \in A(s)} \{U(a(s), m)\}$$

- It follows that

$$U(s, t) = \begin{cases} u(s), s \in \mathcal{T} \\ \min \{U(s', M), s' \in S(s)\} & t = m \\ \max \{U(s', m), s' \in S(s)\}, & t = M \end{cases}$$

Game Strategies: Min-Max Example

- Using the min-max strategy, we compute U as follows:



- Let s be some state and define two quantities:
 - α_s : the minimum utility guaranteed at s thus far.
 - β_s : the maximum utility guaranteed to s thus far.
- If M is the turn-taker at s :
 - α_s increases to the maximum value of s 's successors as they explored.
 - β_s will remain fixed.
- If m is the turn-taker at s :
 - α_s will remain fixed.
 - β_s decreases to the minimum value of s 's successors are explored.
- After exploring a successor of s , if α_s increases beyond β_s , then we need explore any other successors of s since m will ensure s is never reached.

Game Strategies: Expectation-Maximization

- The min-max strategy maximizes M 's minimum utility.
- This means that M will never do worse than what is guaranteed by the min-max strategy, even if m does not always play its best response.
- However, it is very much possible for M to have done better.
- In the **expectation-maximization** strategy, we assume m chooses its action based on a probability distribution, P_s , over $A(s)$.
- It follows that $a_m(s)$, and consequently $U(s, t)$, are random variables.
- We also assume that M wishes to maximize $\mathbb{E}(U(s_0, M))$, where

$$\mathbb{E}[U(s, t)] = \begin{cases} u(s), s \in \mathcal{T} \\ \mathbb{E}[U(a_m(s), M)], t = m \\ \max \{ \mathbb{E}[U(s', m)], s' \in S(s) \}, t = M \end{cases}$$

and $\mathbb{E}[U(a_m(s), M)] = \sum_{a \in A(s)} U(a, M) p_s(a)$.

Game Strategies: Expectation-Maximization Example

- Given p_S as shown below, it seems to make more sense for M to play the right action since there is an 85.5% chance that the resulting utility will be 20.
- However, if M uses the min-max strategy, it would play the action on the left, resulting in an expected utility of ≈ 7.04 .

