

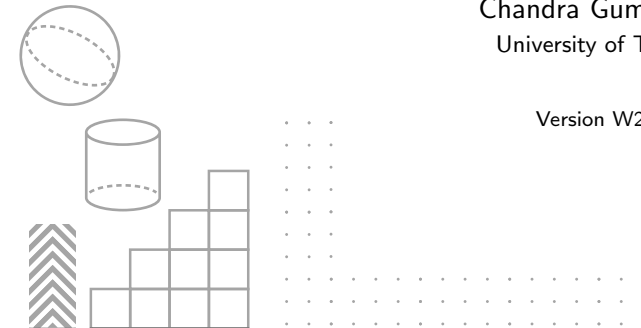
Chapter 10

Knowledge: Reasoning

Introduction to Artificial Intelligence

Chandra Gummaluru
University of Toronto

Version W22.1



- The following is based on material developed by many individuals, including (but not limited to):
 - Sheila McIlraith
 - Bahar Aameri
 - Fahiem Bacchus
 - Sonya Allin

- Resolution assumes the knowledge base is a “clausal theory”.
- A **clausal theory** is a conjunction of “clauses”:
 - If c_1, \dots, c_n are clauses, then $c_1 \wedge \dots \wedge c_n$ is a clausal theory.
- A **clause** is a disjunction of “literals”:
 - If l_1, \dots, l_n are literals, then $l_1 \vee \dots \vee l_n$ is a clause.
- A **literal** is an atomic formula, f , or its negation, $\neg f$.

- Resolution only uses one inference rule:
 - If $c_1 \vee c$ and $c_2 \vee \neg c$ are two clauses, then $c_1 \vee c_2$.
 - The logic here is that c and $\neg c$ cannot simultaneously hold. Therefore, at least one of c_1 or c_2 must hold.

- **Example:** Resolution by Refutation

Consider the following knowledge base:

- ① Clyde is an elephant or a giraffe: $\text{elephant}(\mathbf{Clyde}) \vee \text{giraffe}(\mathbf{Clyde})$.
- ② Either Clyde is not an elephant or he likes peanuts:
 $\neg \text{elephant}(\mathbf{Clyde}) \vee \text{likes}(\mathbf{peanuts}, \mathbf{Clyde})$.
- ③ Either Clyde is not an giraffe or he likes leaves:
 $\neg \text{giraffe}(\mathbf{Clyde}) \vee \text{likes}(\mathbf{leaves}, \mathbf{Clyde})$.
- ④ Clyde does not like leaves: $\neg \text{likes}(\mathbf{leaves}, \mathbf{Clyde})$.

We want to show that Clyde is an elephant. Thus, we assume:

- ⑤ Clyde is not an elephant: $\neg \text{elephant}(\mathbf{Clyde})$

We perform resolution as follows:

- ① elephant(**Clyde**) \vee giraffe(**Clyde**)
- ② \neg elephant(**Clyde**) \vee likes(**peanuts**, clyde)
- ③ \neg giraffe(**Clyde**) \vee likes(**leaves**, **Clyde**)
- ④ \neg likes(**leaves**, **Clyde**)
- ⑤ \neg elephant(**Clyde**)
- ⑥ R[5a,1a] giraffe(**Clyde**)
- ⑦ R[6a,3a] likes(**leaves**, **Clyde**)
- ⑧ R[7a,4a] {}

- Often, assertions in our knowledge base are not expressed in clausal form.

Example: Non-Clausal Assertion

- If Clyde is an elephant, then he likes peanuts:
elephant(**Clyde**) \rightarrow likes(**peanuts**, **Clyde**).
- We need a systemic way to convert any non-clausal assertion into a set of clauses.

- The following procedure can be used:

- ① Eliminate \rightarrow
- ② Move \neg inward
- ③ Standardize variables
- ④ Eliminate \exists
- ⑤ Move \forall outward
- ⑥ Distribute \vee over \wedge
- ⑦ Flatten nested \wedge/\vee .
- ⑧ Remove \forall
- ⑨ Split on \wedge

① Eliminate \rightarrow

- $A \rightarrow B \equiv \neg A \vee B$

E.g: If Clyde is an elephant, he likes peanuts. Equivalently, either Clyde likes peanuts, or he is not an elephant.

② Move \neg inward and simplify:

- $\neg\neg A \equiv A$

E.g: Clyde isn't not an elephant iff he's an elephant.

- $\neg(A \vee B) \equiv \neg A \wedge \neg B$

E.g: Clyde is neither a tiger nor a giraffe iff he isn't a tiger and he isn't a giraffe.

- $\neg(A \wedge B) \equiv \neg A \vee \neg B$

E.g: Clyde doesn't like both leaves and meat iff he dislikes leaves or meat.

- $\neg\forall x A \equiv \exists x \neg A$

E.g: Not every person likes this class iff some people don't like it.

- $\neg\exists x A \equiv \forall x \neg A$

E.g: Not one person dislikes this class iff everyone likes this class.

③ Distinguish quantified variables:

- $\forall x f(x) \equiv \forall y f(y)$
- $\exists x f(x) \equiv \exists y f(y)$

④ Eliminate \exists :

- $\exists x[f(x)] \equiv f(\mathbf{c})$, for some unique constant, \mathbf{c} .
E.g: If there is a friendly elephant, call him “Clyde”. If “Clyde” is a friendly elephant, then there exists a friendly elephant.
- $\forall x[\exists yf(y)] \equiv \forall x[f(g(x))]$, for some g .
E.g: Everyone likes someone (usually a different someone). If we use $\text{partnerOf}(\cdot)$ to denote that person, we can say, everyone likes their partner.

⑤ Distribute \vee over \wedge :

- $A \vee (B \wedge C) \equiv (A \vee B) \wedge (A \vee C)$.

E.g: Clyde is either a giraffe, or he is an elephant and likes peanuts iff Clyde is either a giraffe or an elephant, and he is either a giraffe or likes peanuts.

⑥ Flatten nested \wedge/\vee :

- $(A \wedge (B \wedge C)) \equiv (A \wedge B \wedge C)$
- $(A \vee (B \vee C)) \equiv (A \vee B \vee C)$

⑦ Remove \forall :

- $\forall x f(x) \equiv f(x)$

Converting an Assertion to Clausal Form: Example

- Consider the statement:

$$\forall x \left[P(x) \rightarrow \left(\left(\forall y [P(y) \rightarrow P(f(x, y))] \right) \wedge \neg \left(\forall y [\neg q(x, y) \wedge P(y)] \right) \right) \right]$$

- We can convert it to clausal form as follows:

- 1 Eliminate \rightarrow

$$\forall x \left[\neg P(x) \vee \left(\left(\forall y [\neg P(y) \vee P(f(x, y))] \right) \wedge \neg \left(\forall y [\neg q(x, y) \wedge P(y)] \right) \right) \right]$$

- 2 Move \neg inward

$$\forall x \left[\neg P(x) \vee \left(\left(\forall y [\neg P(y) \vee P(f(x, y))] \right) \wedge \left(\exists y [q(x, y) \vee \neg P(y)] \right) \right) \right]$$

Converting an Assertion to Clausal Form: Example

- ② Move \neg inward

$$\forall x \left[\neg P(x) \vee \left(\left(\forall y [\neg P(y) \vee P(f(x, y))] \right) \wedge \left(\exists y [q(x, y) \vee \neg P(y)] \right) \right) \right]$$

- ③ Standardize variables

$$\forall x \left[\neg P(x) \vee \left(\left(\forall y [\neg P(y) \vee P(f(x, y))] \right) \wedge \left(\exists z [q(x, z) \vee \neg P(z)] \right) \right) \right]$$

- ④ Eliminate \exists

$$\forall x \left[\neg P(x) \vee \left(\left(\forall y [\neg P(y) \vee P(f(x, y))] \right) \wedge \left([q(x, g(x)) \vee \neg P(g(x))] \right) \right) \right]$$

- ⑤ Move \forall outward

$$\forall x \forall y \left[\neg P(x) \vee \left(\left(\neg P(y) \vee P(f(x, y)) \right) \wedge \left([q(x, g(x)) \vee \neg P(g(x))] \right) \right) \right]$$

Converting an Assertion to Clausal Form: Example

- 5 Move \forall outward

$$\forall x \forall y \left[\neg P(x) \vee \left(\left(\neg P(y) \vee P(f(x, y)) \right) \wedge \left([q(x, g(x)) \vee \neg P(g(x))] \right) \right) \right]$$

- 6 Distribute \vee over \wedge

$$\forall x \forall y \left[\left(\neg P(x) \vee \left(\neg P(y) \vee P(f(x, y)) \right) \right) \wedge \left(\neg P(x) \vee [q(x, g(x)) \vee \neg P(g(x))] \right) \right]$$

- 7 Flatten nested \wedge/\vee

$$\forall x \forall y \left[\left(\neg P(x) \vee \neg P(y) \vee P(f(x, y)) \right) \wedge \left(\neg P(x) \vee q(x, g(x)) \vee \neg P(g(x)) \right) \right]$$

Converting an Assertion to Clausal Form: Example

- 7 Flatten nested \wedge/\vee

$$\forall x \forall y \left[\left(\neg P(x) \vee \neg P(y) \vee P(f(x, y)) \right) \wedge \left(\neg P(x) \vee q(x, g(x)) \vee \neg P(g(x)) \right) \right]$$

- 8 Remove \forall

$$\left(\neg P(x) \vee \neg P(y) \vee P(f(x, y)) \right) \wedge \left(\neg P(x) \vee q(x, g(x)) \vee \neg P(g(x)) \right)$$

- 9 Split on \wedge

$$\begin{cases} \neg P(x) \vee \neg P(y) \vee P(f(x, y)) \\ \neg P(x) \vee q(x, g(x)) \vee \neg P(g(x)) \end{cases}$$

- In the previous example of resolution by refutation, none of the conflicting clauses had variables.
- Suppose instead that we had two conflicting clauses but at least one involves some variables such as $(p(x) \vee p'(y) \dots)$ and $(\neg p(\mathbf{a}) \vee \dots)$.
 - Since x can be any constant, the clause $(p(x) \vee p'(y) \vee \dots)$ actually represents a family of clauses

$$(p(\mathbf{a}) \vee p'(y) \vee \dots)$$

$$\vdots$$

$$(p(\mathbf{z}) \vee p'(y) \vee \dots)$$

- Thus, we can still resolve the clauses by substituting $x = \mathbf{a}$, yielding $(p'(y) \vee \dots)$.
- We could have also made additional substitutions, like $y = \mathbf{b}$, however, this would reduce the generality of the resulting clause.

Unification of Clauses: Substitutions

- A **substitution** is a finite set of equations of the form, $V = t$, where V is a variable, and t is a term not containing V .
 - Applying the substitution, $\delta = \{V_1 = t_1, \dots, V_n = t_n\}$, to the formula f , is done by simultaneously replacing V_i with t_i .
 - The resulting formula is denoted $f\delta$.
 - **E.g:** If $f = P(x, g(y, z))$, and $\delta = \{x = y, y = f(a)\}$, then $f\delta = P(y, g(f(a), z))$.
- We can compose two substitutions, θ, δ to obtain a new substitution, $\theta\delta$.
 - Let $\theta = \{x_1 = s_1, \dots, x_n = s_n\}$ and $\delta = \{y_1 = t_1, \dots, y_m = t_m\}$.
 - We compute $\theta\delta$ as follows:
 - ① For each equation, $x_i = s_i$ in θ , apply δ to its right side to yield $x_i = s_i\delta$.
 - ② If $x_i = s_i\delta$ is not a tautology (e.g., $V = V$), include it in $\theta\delta$.
 - ③ If $y_i \neq x_j$ for any j , then include $y_i = t_i$ in $\theta\delta$.
 - Defining composition in this way means that applying $\theta\delta$ to a formula is equivalent to first applying θ , and then applying δ , i.e., $(f\theta)\delta = f(\theta\delta)$.

- **Example:** Composing Substitutions

- Let $\theta = \{x = f(y), y = z\}$ and $\delta = \{x = a, y = b, z = y\}$.
 - ① Applying δ to the right side of $x = f(y)$ and $y = z$ yields $x = f(b)$ and $y = y$, respectively.
 - ② Since $y = y$ is a tautology, we do not include it.
 - ③ We also include $z = y$ from δ .
 - ④ It follows that $\theta\delta = \{x = f(b), z = y\}$.

- A **unifier** of two formulae, f and g , is a substitution, δ , that makes f and g syntactically identical.
- A unifier, δ , is **most general** iff for every other unifier, θ , there exists a third substitution, λ , such that

$$\theta = \delta\lambda.$$

In other words, every other unifier is more specialized.

Unification of Clauses: Computing the Most General Unifier

- **Procedure:** Computing the Most General Unifier
 - Let f, g be two formulae we wish to unify.
 - Let δ denote the most general unifier, and $S = \{f, g\} \delta$.
- ① Start with the empty substitution, $\delta_0 = \{\}$, and $S_0 = \{f, g\}$
- ② In each iteration, k , find a disagreement set, $D_k = \{V, t\}$. If one does not exist, then $\delta = \delta_k$ is the most general unifier, and $S = S_k$.
- ③ If D_k is unifiable (i.e., V is a variable, and t is a term not containing V):
 - a Let $\delta_{k+1} = \delta_k \{V = t\}$.
 - b Let $S_{k+1} = S_k \{V = t\}$.
- ④ If D_k is not unifiable, then f and g cannot be unified.

Computing the Most General Unifier: Example 1

- In this example, we consider finding the most general unifier of $P(\mathbf{a}, x, h(g(z)))$ and $P(z, h(z), h(y))$.

ltr. D_k	S_k	δ_k
0 $\{\mathbf{a}, z\}$	$\left\{ P(\mathbf{a}, x, h(g(z))), P(z, h(y), h(y)) \right\}$	$\{\}$
1 $\{x, h(y)\}$	$\left\{ P(\mathbf{a}, x, h(g(z))), P(\mathbf{a}, h(y), h(y)) \right\}$	$\{z = \mathbf{a}\}$
2 $\{x, h(y)\}$	$\left\{ P(\mathbf{a}, h(y), h(g(\mathbf{a}))), P(\mathbf{a}, h(y), h(y)) \right\}$	$\{z = \mathbf{a}, x = h(y)\}$
3 $\{h(g(\mathbf{a})), h(y)\}$	$\left\{ P(\mathbf{a}, h(g(\mathbf{a})), h(g(\mathbf{a}))) \right\}$	$\{z = \mathbf{a}, x = h(g(\mathbf{a})), y = g(\mathbf{a})\}$

- The most general unifier is $\delta = \{z = \mathbf{a}, x = h(g(\mathbf{a})), y = g(\mathbf{a})\}$.

Computing the Most General Unifier: Example 2

- In this example, we consider finding the most general unifier of $P(f(\mathbf{a}), g(x))$ and $P(y, y)$.

	ltr. D_k	S_k	δ_k
0	$\{f(\mathbf{a}), y\}$	$\{P(f(\mathbf{a}), g(x)), P(y, y)\}$	$\{\}$
1	$\{g(x), f(\mathbf{a})\}$	$\{P(f(\mathbf{a}), g(x)), P(f(\mathbf{a}), f(\mathbf{a}))\}$	$\{y = f(\mathbf{a})\}$

- The disagreement set, D_1 is not unifiable. Thus, there is no unifier.

- At U of T, it is reasonable to assume that:
 - i Some courses have exams.
 - ii No student likes anything with an exam.
- Can we show that no student likes all courses?

① Pick a vocabulary to represent the assertions:

- $\text{isStudent}(x)$, x is a student
- $\text{isCourse}(x)$, x is a course
- $\text{hasExam}(x)$, x has an exam
- $\text{likes}(x, y)$, y likes x

- ② Convert each assertion (including the negation of the query) to a first-order formula:

- i Some courses have exams:

$$\exists x[\text{isCourse}(x) \wedge \text{hasExam}(x)]$$

- ii No student likes anything that has an exam:

$$\forall x \forall y [\text{isStudent}(x) \wedge \text{hasExam}(y) \rightarrow \neg \text{likes}(y, x)]$$

- iii Some student likes all courses:

$$\exists x [\text{isStudent}(x) \wedge \forall y [\text{isCourse}(y) \rightarrow \text{likes}(y, x)]]$$

③ Convert each first-order formula to clausal form:

i $\exists x[\text{isCourse}(x) \wedge \text{hasExam}(x)]:$

- C1F.4: $\text{isCourse}(c) \wedge \text{hasExam}(c)$
- C1F.8: $\text{isCourse}(c), \text{hasExam}(c)$

ii $\forall x \forall y[\text{isStudent}(x) \wedge \text{hasExam}(y) \rightarrow \neg \text{likes}(x, y)]:$

- C1F.2: $\forall x \forall y [\neg (\text{isStudent}(x) \wedge \text{hasExam}(y)) \vee \neg \text{likes}(x, y)]$
- C1F.3: $\forall x \forall y [\neg \text{isStudent}(x) \vee \neg \text{hasExam}(y) \vee \neg \text{likes}(x, y)]$
- C1F.8: $\neg \text{isStudent}(x) \vee \neg \text{hasExam}(y) \vee \neg \text{likes}(x, y)$

iii $\exists x[\text{isCourse}(x) \wedge \forall y[\text{isStudent}(y) \rightarrow \text{likes}(x, y)]]:$

- C1F.2: $\exists x[\text{isStudent}(x) \wedge \forall y[\neg \text{isCourse}(x) \vee \text{likes}(x, y)]]$
- C1F.4: $\text{isStudent}(s) \wedge \forall y[\neg \text{isCourse}(y) \vee \text{likes}(s, y)]$
- C1F.5: $\forall y[\text{isStudent}(s) \wedge [\neg \text{isCourse}(y) \vee \text{likes}(s, y)]]$
- C1F.8: $\text{isStudent}(s), (\neg \text{isCourse}(y) \vee \text{likes}(s, y))$

④ Apply the inference rule of resolution to the clauses:

- ① $\text{isCourse}(\mathbf{c})$
- ② $\text{hasExam}(\mathbf{c})$
- ③ $(\neg \text{isStudent}(x) \vee \neg \text{hasExam}(y) \vee \neg \text{likes}(x, y))$
- ④ $\text{isStudent}(\mathbf{s})$
- ⑤ $\neg \text{isCourse}(z), \text{likes}(\mathbf{s}, z)$

- ⑥ $R[1a, 5a] \{z = \mathbf{c}\} \text{likes}(\mathbf{s}, \mathbf{c})$
- ⑦ $R[6a, 3c] \{x = \mathbf{s}, y = \mathbf{c}\} \neg \text{isStudent}(\mathbf{s}) \vee \neg \text{hasExam}(\mathbf{c})$
- ⑧ $R[7a, 4a] \neg \text{hasExam}(\mathbf{c})$
- ⑨ $R[8a, 2a] \{\}$

Resolution by Refutation: Answer Extraction

- So far, we can answer any query where the answer is yes/no.
- Queries often are more nuanced.
 - **E.g:** Harry and Ron are friends and Harry and Ginny are married. Assuming all married couples are friends, who are Harry's friends?
- To answer such a query, we can assume the answer, which we denote using an answer clause, $\text{answer}(x)$.
 - Either x is the answer or x and Harry and x are not friends:

$$\neg \text{areFriends}(\mathbf{Harry}, x) \vee \text{answer}(x)$$

- We then perform resolution until we obtain a clause of only answer literals.

- We could perform resolution as follows:
 - ① $\text{areFriends}(\mathbf{Harry}, \mathbf{Ron})$
 - ② $\text{areMarried}(\mathbf{Harry}, \mathbf{Ginny})$
 - ③ $\neg \text{areMarried}(x, y) \vee \text{areFriends}(x, y)$
 - ④ $\neg \text{areFriends}(\mathbf{Harry}, x) \vee \text{answer}(x)$
 - ⑤ $R[1a, 4a] \{x = \mathbf{Ron}\} \text{answer}(\mathbf{Ron})$
 - ⑥ $R[2a, 3a] \{x = \mathbf{Harry}, y = \mathbf{Ginny}\} \text{areFriends}(\mathbf{Harry}, \mathbf{Ginny})$
 - ⑦ $R[6a, 4a] \{x = \mathbf{Ginny}\} \text{answer}(\mathbf{Ginny})$
- Ron and Ginny are Harry's friends.