



Games

Formalization and Algorithms

Introduction to Artificial Intelligence

Chandra Gummaluru

University of Toronto

Version W22.1

- The following is based on material developed by many individuals, including (but not limited to):
 - Sheila McIlraith
 - Bahar Aameri
 - Fahiem Bacchus
 - Sonya Allin

- Recall the components of a **game**:
 - N players, indexed 1 through N
 - a set of states, S , in which the game could be in
 - a set of terminal states, $T \subseteq S$ in which the game ends
 - a set of actions, $A_i(s)$, available to each player, i from each state $s \in S$:
 - each n -tuple, (a_1, \dots, a_N) , $a_i \in A_i(s)$ is called a joint-action from s , and $A(s)$ denotes the set of all possible joint actions, i.e.,

$$A(s) = \prod_{j=1}^N A_j(s).$$

- we use a_{-i} to denote the joint action between all players except i , and $A_{-i}(s)$ to denote the set of all possible joint actions between all players except i , i.e.,

$$A_{-i}(s) = \prod_{j \neq i} A_j(s)$$

- we use $S(s, a)$ to denote the state resulting from applying $a \in A(s)$ to s
- a reward function, $r_i : A \rightarrow \mathbb{R}$, for each player i , so that $r_i(a)$ is a measure of how useful an action a is to player i

- Given an $s_0 \in S$, each possible realization of the game is a sequence of actions, $\langle a^{(1)}, \dots, a^{(n)} \rangle$, such that $a^{(k)} \in A(s_{k-1})$, $s_k = S(s_{k-1}, a^{(k)})$ and $s_n \in T$.
- The dynamics of the game are determined by a policy function, p , where $p(a|s)$ is the probability that $a \in A(s)$ is played from state, s .
- To solve a game means to find p assuming the players are rational.
- We can define the the expected-reward (value) of taking the action, a_i , for player i , when the other players' joint action is a_{-i} , using the following recurrence:

$$v_i(a_i, a_{-i}) = \begin{cases} r_i(a_i, a_{-i}) & S(s, a_i, a_{-i}) \in T \\ r_i(a_i, a_{-i}) + \sum_{a' \in A(s')} p(a'|s') v_i(a') & \text{otherwise} \end{cases} \quad (\text{ExpRwd})$$

where $a_i \in A_i(a)$, $a_{-i} \in A_{-i}(s)$, and $s' = S(s, a)$.

Formalizing 2-Player Sequential Zero-Sum Deterministic Games

- If all players are rational, we would have

$$p(a|s) = \begin{cases} 1 & a = a^*(s) \\ 0 & \text{otherwise} \end{cases} \quad (\text{OptPol})$$

where $a^*(s)_i$ is player i 's best-response to the joint action, $a^*(s)_{-i}$:


$$a^*(s)_i = \arg \max_{a_i \in A_i(s)} v_i(a_i, a^*(s)_{-i}). \quad (\text{OptAct})$$

- It is difficult to find p in general, so we made two simplifying assumptions:
 - There is only one-player.
 - The set of terminal states, $T \subseteq S$, is replaced with a set of goal states, $G \subseteq S$
- We shall now slightly relax both of these assumptions:
 - there are now two-players
 - only one player, called the turn-taker, can act from any given state; rather than defining A_1 and A_2 separately, we simply define $A(s)$ to be the set of actions available to the turn-taker at s
 - the turn-taker switches after each action
 - we assume $r_1(a) = -r_2(a)$; rather than defining r_1 and r_2 separately, we let $r(a)$ to be the reward obtained by the turn-taker for taking action, a , and assume that r_1
 - $r(a) = 0$ for any $a \in A(s)$ such that $S(s, a) \notin T$

Formalizing a Search Problem: Search Trees

- We saw that a graph could be used to represent the structure of a search problem. However, the searching itself takes place on a tree of possible paths.


Example: Search Tree for Slider Puzzle



Formalizing a Game: Game Trees

- Similarly, we can create a tree of possible paths for our game. However in this case, each level corresponds to different player, namely, the turn-taker.

Example: Game Tree for Tic-Tac-Toe Puzzle



Formalizing 2-Player Sequential Zero-Sum Deterministic Games


- Under the new assumptions, (ExpRwd) becomes:

$$v(a) = \begin{cases} r(a) & S(s, a) \in T \\ - \sum_{a' \in A(s')} p(a'|s) v(a') & \text{otherwise} \end{cases}$$

- The strategy of a player is defined by the family of distributions, $p(\cdot|s)$ for s in which they are the turn-taker.

Game Strategies: Example

- We will analyze various strategies on the tree below, where \bullet and \circ respectively denote states in which we are the turn-taker, or our adversary is:



- The rewards are expressed in the perspective of the player who last acted.

- If both players are perfect, i.e., p satisfies (OptPol), then (ExpRwd) reduces to:

$$v(a) = \begin{cases} r(a) & S(s, a) \in T \\ - \max_{a' \in A(s')} v(a') & \text{otherwise} \end{cases} \quad (\text{NegMax})$$

where $s' = S(s, a)$.


- We can equivalently, express (NegMax) in terms of v_1 only, in which case, we have

$$v_1(a) = \begin{cases} r_1(a) & S(s, a) \in T \\ \max_{a' \in A(s')} v_1(a') & S(s, a) \notin T \text{ and we are the turn-taker} \\ \min_{a' \in A(s')} v_1(a') & S(s, a) \notin T \text{ and we are not the turn-taker} \end{cases} \quad (\text{MinMax})$$


- We can define the value of states as follows: $v_1(S(s, a)) = v_1(a), \forall a \in A(s), s \in S$

Game Strategies: Perfect Players

- To use (MinMax), we need to first compute the rewards in our perspective.



- We can now compute the rewards of the actions as follows:




Game Strategies: MinMax Psuedocode

- For any non-terminal state, we recursively compute either the maximum or minimum of its successors' values.

```
1: procedure MINMAXSEARCH( $s, t$ )
2:   if  $s \in \mathcal{T}$  then                                 $\triangleright s$  is a terminal state
3:     return  $r_1(s)$ 
4:   if  $t = 1$  then
5:      $v_1(s) \leftarrow -\infty$ 
6:     for  $a \in A(s)$  do
7:        $v_1(s) \leftarrow \max \{v_1(s), \text{MINMAXSEARCH}(S(s, a), 2)\}$ 
8:   if  $t = 2$  then
9:      $v_1(s) \leftarrow \infty$ 
10:    for  $a \in A(s)$  do
11:       $v_1(s) \leftarrow \min \{v_1(s), \text{MINMAXSEARCH}(S(s, a), 1)\}$ 
12:    return  $v_1(s)$ 
```


- It turns out if we use a depth-first search, we do not necessarily need to explore the entire tree to compute the min-max utility of the root.
- Let s be some state and define two quantities:
 - α_s : the maximum utility guaranteed at s thus far.
 - β_s : the minimum utility guaranteed at s thus far.

- If we are the turn-taker at s :
 - α_s increases to the maximum value of s 's successors as they explored.
 - $\beta_s = \beta_{pt(s)}$ and will remain fixed.



- After exploring a successor of s , if α_s increases beyond β_s , then we need not explore any other successors of s since our adversary will ensure s is never reached.


- If our adversary is the turn-taker at s :
 - $\alpha_s = \alpha_{\text{pt}(s)}$ and will remain fixed.
 - β_s decreases to the minimum value of s 's successors as they explored.



- After exploring a successor of s , if β_s decreases beyond α_s , then we need not explore any other successors of s since we will ensure s is never reached.

Game Strategies: Minimax w/ α/β -Pruning Example

- Using the min-max strategy with α/β -pruning, we prune the nodes shown below:



Game Strategies: Min-Max w/ α/β -Pruning Psuedocode

```
1: procedure ALPHABETASEARCH( $s, t, \alpha, \beta$ )
2:   if  $s \in T$  then                                ▷  $s$  is a terminal state
3:     return  $r_1(s)$ 
4:   if  $t = 1$  then
5:     for  $a \in A(s)$  do
6:        $\alpha \leftarrow \max \{ \alpha, \text{ALPHABETASEARCH}(S(s, a), 2, \alpha, \beta) \}$ 
7:       if  $\alpha \geq \beta$  then
8:         break
9:     return  $\alpha$ 
10:  if  $t = 2$  then
11:    for  $a \in A(s)$  do
12:       $\beta \leftarrow \min \{ \alpha, \text{ALPHABETASEARCH}(S(s, a), 1, \alpha, \beta) \}$ 
13:      if  $\alpha \geq \beta$  then
14:        break
15:    return  $\beta$ 
```

- The (MinMax) strategy assumes that our adversary always plays perfectly; thus, it maximizes our minimum expected reward.
- This is why it is a good strategy to use if we do not have any other knowledge of our adversary.
- However, in most cases, we have some estimate of $p(\cdot|s)$, whenever the adversary is the turn-taker at s .
- In such cases, it is very much possible for us to have done better.

- In this strategy, we keep an estimate of our adversary's strategy, $\tilde{p}(\cdot|s)$.
- We will study how to compute \tilde{p} later.
- For now, we simply use it in (ExpRwd) and get

$$v_1(a) = \begin{cases} r_1(a) & S(s, a) \in T \\ \max_{a' \in A(s')} v_1(a') & S(s, a) \notin T \text{ and we are the turn-taker} \\ \sum_{a' \in A(s')} v_1(a') \tilde{p}(a'|s) & S(s, a) \notin T \text{ and we are not the turn-taker} \end{cases}$$

(ExpMax)

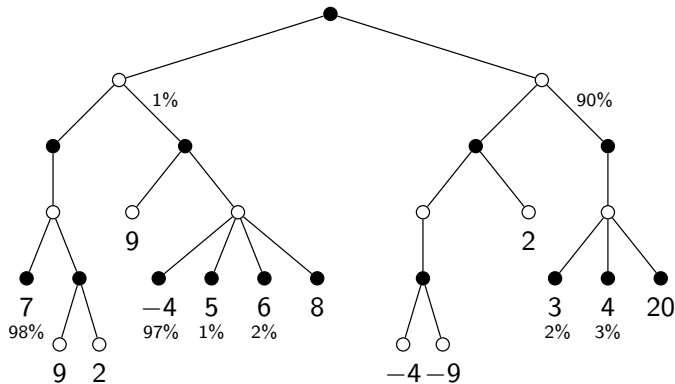
Game Strategies: ExpectMax Psuedocode

- For any non-terminal circumstance, we recursively compute either the maximum or expected-value of its successors' values.

```
1: procedure EXPECTMAXSEARCH( $s, t$ )  
2:   if  $s \in \mathcal{T}$  then                                ▷  $s$  is a terminal state  
3:     return  $r_1(s)$   
4:   if  $t = 1$  then  
5:      $v_1(s) \leftarrow -\infty$   
6:     for  $a \in A(s)$  do  
7:        $v_1(s) \leftarrow \max \{v_1(s), \text{EXPECTMAXSEARCH}(S(s, a), 2)\}$   
8:   if  $t = 2$  then  
9:      $v_1(s) \leftarrow 0$   
10:    for  $a \in A(s)$  do  
11:       $v_1(s) \leftarrow v_1(s) + \tilde{p}(a|s) \times \text{EXPECTMAXSEARCH}(S(s, a), 1)$   
12:    return  $v_1(s)$ 
```


Game Strategies: Expectimax Example

- Suppose $p(\cdot|s)$ was known for all s in which the adversary is the turn-taker:



Game Strategies: Expectimax Example

- We can compute $p(\cdot|s_0)$ as follows:



- A fundamental problem with our approach is that we must explore the entire game-tree before making a decision.
- If we do not have time to explore the entire tree, we need some way to estimate the state values (i.e., a heuristic).
- We can use this estimate in place of the actual expected rewards.

Example: Heuristic for Tic-Tac-Toe Puzzle

- Count the number of ways we could win via our moves thus far.

