



Search

# Informed Algorithms

Introduction to Artificial Intelligence


Chandra Gummaluru

University of Toronto

Version W22.1

- The following is based on material developed by many individuals, including (but not limited to):
  - Sheila McIlraith
  - Bahar Aameri
  - Fahiem Bacchus
  - Sonya Allin

- An ideal search algorithm would explore “more promising” paths first.
- **Example:** Informed Search in a Superstore
  - You enter a super-store and want a T-shirt. Where would you look first?




- You would probably check the “Apparel” section even though it is farther away.

- We want to develop a metric to guide the order in which paths are explored.
- We will represent this metric as a function,  $f$ , of the path; paths with smaller  $f$ -values should be explored earlier.
- So far, we have been defining the  $f$ -value of a path in terms of its length/cost.
- However, we now also want to use info about the states traversed by the path. In particular, we want to measure how “good” the final state in the path is.
- Such algorithms are said to be **informed**.

- We define a function,  $h$ , called a **heuristic**, so that  $h(s)$  ideally estimates the minimum cost of getting from a state,  $s$ , to some goal state.
- We extend  $h$  to operate on paths by defining  $h(p) = h(s_n)$  where  $p = \langle s_0, \dots, s_n \rangle$ .
- In other words, the  $h$ -value of a path is the  $h$ -value of its terminal state.
- We can now use  $h$  as part of  $f$ . Two common choices include:
  - ①  $f(p) = h(p)$ , which is called **greedy best-first search** (GBFS)
  - ②  $f(p) = c(p) + h(p)$ , which is called **A-star** ( $A^*$ )


- **Example:** Informed Search in a Superstore
  - You want to minimize the total estimated cost, i.e.,  $f(p) = c(p) + h(p)$ .



- Although  $c(\langle \text{Doors}, \text{Electronics} \rangle), c(\langle \text{Doors}, \text{Produce} \rangle) < c(\langle \text{Doors}, \text{Apparel} \rangle)$ , we check “Apparel” first since we feel  $h(\text{Apparel}) \ll h(\text{Electronics}), h(\text{Produce})$ .

# Toy Search Graph for Informed Search

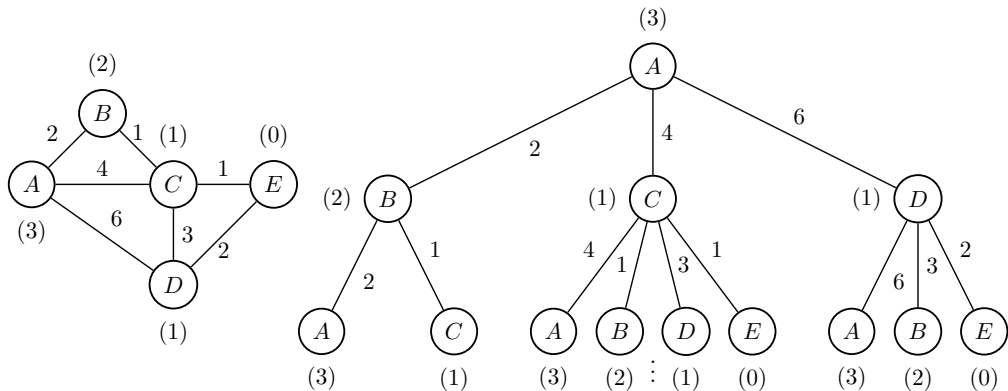
- To compare the effects of different  $f$ -functions, we will run the corresponding algorithms on the following search graph (starting at  $A$  and looking for  $E$ ):



- We will use path-checking and break ties alphabetically.

# Tree of Possible Paths: Example

- Recall that the generalized search algorithm explores a tree of possible paths.




- We wish to determine the properties of the various search algorithm in terms of  $b$ ,  $d$ ,  $\varepsilon$ ,  $m$  and  $c^*$ , where:
  - $b$ ,  $m$ , and  $\varepsilon$ , be the branching factor, depth, and minimum edge weight of the tree
  - $d$  and  $c^*$  denote the length and cost of the optimal solution



## Greedy Best-First Search: Example

- In greedy best-first search (GBFS), we use  $f(p) = h(p)$ . In other words, a partial path is more promising if we estimate that the remaining cost to the goal is lower.

ltr.	$p$	$\mathcal{O}$
1	–	$A 12$
2	$A$	$AB 8, AC 2, AD 2$
3	$AC$	$AD 2, AB 8, ACB 8, ACD 2, ACE 0$
4	$ACE$	$AD 2, AB 8, ACB 8, ACD 2$



- Notice that this particular heuristic yielded a sub-optimal solution, but it did reduce the number iterations needed to find a solution (relative to UCS).


- **Complete** for  $b, d < \infty$  if path-checking is used.
  - There are at most  $1 + b + \dots b^d$  paths in the tree.
  - The search will never explore any path more than once.
  - The search must explore all paths reachable from the root.
- **Optimal** never<sup>1</sup>.
  - $f$  does not incorporate  $c$ .
- **Time Complexity:**  $O(b^m)$ .
  - In the worst-case, we will need to explore every path.
- **Space Complexity:**  $O(bm)$ .
  - In the worst-case, we will need to explore every path.

---

<sup>1</sup>except by chance.

- In A-Star ( $A^*$ ), we use  $f(p) = c(p) + h(p)$ . In other words, a partial path is more promising if we estimate its total cost to the goal to be lower.

ltr.	$p$	$\mathcal{O}$
1	–	$A 12$
2	$A$	$AB 10, AC 6, AD 8$
3	$AC$	$AB 10, AD 8, ACB 13, ACD 9, ACE 5$
4	$ACE$	$AD 8, AB 10, ACB 13, ACD 9$



- Notice that this particular heuristic yielded a sub-optimal solution, but it did reduce the number iterations needed to find a solution (relative to UCS).

- The optimal path from  $A$  to  $E$  is  $\langle A, B, C, E \rangle$ .
- However, both GBFS and  $A^*$  found the same sub-optimal path,  $\langle A, C, E \rangle$ .
- This is expected in GBFS since it does not make use of the costs.
- Since  $A^*$  does make use of the costs, ideally the path it finds would be optimal.
- The problem was that the heuristic we used over-estimated the cost to get from  $B$  to  $E$ , making the partial path  $\langle A, B \rangle$  seem far less promising.

- If  $h^*(s)$  is the true cost from  $s$  to the nearest goal, then ideally, our heuristic would be such that  $h(s) \leq h^*(s)$  for all  $s$ , i.e.,  $h$  never over-estimates the true cost.
- Such a heuristic is said to be **admissible**.
- If a heuristic does not over-estimate the cost of individual actions, i.e.,

$$h(s) - h(a(s)) \leq c(a), \forall s, a \in A(s),$$

we say that the heuristic is **consistent** or **monotone**.

- Any consistent heuristic,  $h$ , where  $h(s) = 0$  for every  $s \in G$ , is also admissible.

- Proof:**

- Pick an arbitrary initial state,  $s_0$ .
- If no goal-terminating path from  $s_0$  exists, then  $h^*(s_0)$  is infinite and the claim trivially holds.
- Otherwise, let  $\langle s_0, \dots, s_n \rangle$  be resulting state sequence of the optimal goal-terminating path,  $\langle a_{0,1}, \dots, a_{n-1,n} \rangle$  from  $s_0$ .
- Since  $s_n \in G$ , it follows that  $h(s_n) = 0 = h^*(s_n)$ . Thus,  $h(s_n) \leq h^*(s_n)$ .
- Assume  $h(s_{i+1}) \leq h^*(s_{i+1})$  for some  $i$ . Then,  $h(s_i) \leq h^*(s_i)$ :

$$h(s_i) \leq c(a_{i,i+1}) + h(s_{i+1}) \leq c(a_{i,i+1}) + h^*(s_{i+1}) = h^*(s_i).$$

- Since  $s_0$  was arbitrary, it follows that  $h(s) \leq h^*(s)$  for all  $s$ , i.e.,  $h$  is admissible.

- It turns out that if  $h$  is consistent, then A-star satisfies the condition needed to use global path checking.
- We prove this in three parts, all of which rely on  $h$  being consistent:
  - ① If  $f \equiv c + h$ , then  $f$  must be non-decreasing along any path.
  - ② The  $f$ -values of paths explored in A\* are non-decreasing.
  - ③ The first time A\* finds a path to a state, it has found the optimal path.

① If  $f \equiv c + h$ , then the  $f$ -value must be non-decreasing along any path.

- Let  $(s_0, \langle a_{0,1}, \dots, a_{n-1,n} \rangle)$  be a path, and  $\langle a_0, \dots, s_n \rangle$ .
- By the definition of  $f$  in  $A^*$ ,

$$f(s_i) = c(a_{0,1}, \dots, a_{i-1,i}) + h(s_i).$$

- If  $h$  is consistent, then by definition,  $h(s_i) \leq c(a_{i,i+1}) + h(s_{i+1})$ .
- It follows that  $f(s_i) \leq c(a_{0,1}, \dots, a_{i,i+1}) + h(s_{i+1}) = f(s_{i+1})$



- ② The  $f$ -values of paths explored in A\* are non-decreasing.
- If  $p'$  is explored after  $p$ , then either:
    - $p'$  was already on the open when  $p$  was explored
    - $p'$  is an extension of  $p$
    - $p'$  is an extension of a path,  $p''$  such that  $f(p'') \geq f(p)$ .
  - In any case  $f(p') \geq f(p)$ .


- ③ The first time A\* finds a path to a state, it has found the optimal path.
- Let  $p = \langle s_0, s_1, \dots, s_{n-1}, s_n \rangle$  and  $p' = \langle s_0, s'_1, \dots, s'_{n-1}, s_n \rangle$  be two state sequences to the goal state,  $s_n$  found by A\*.
  - If  $p$  was found before  $p'$ , then by (2),  $f(p') \geq f(p)$ .
  - Using the definition of  $f$  in A\*, we have

$$c(p') + h(p') \geq c(p) + h(p).$$

- Since  $h$  is consistent (so admissible),  $h(p') = h(p) = 0$ , and so  $c(p') \geq c(p)$ .

- Let us reconsider A\* using a consistent heuristic.

ltr.	$p$	$\mathcal{O}$
1	–	$A 3$
2	$A$	$AB 4, AC 5, AD 7$
3	$AB$	$AC 5, AD 7, ABC 4$
4	$ABC$	$AC 5, AD 7, ABCE 5, ABCD 7$
5	$ABCE$	$AC 5, AD 7, ABCD 7$



- Notice that this particular heuristic yielded the optimal solution.

- **Complete** for  $b, d < \infty, \varepsilon > 0$ .
  - There are at most  $1 + b + \dots b^d$  paths in the tree.
  - The search will never explore any path more than once.
  - The search must explore all paths reachable from the root.
- **Optimal** if  $h(\cdot)$  is admissible.
  - Let  $s_n$  be a goal state and  $p = \langle s_0, \dots, s_n \rangle$  be the path found by A\*.
  - Recall that for A\*,  $f(\cdot) = c(\cdot) + h(\cdot)$ .
  - If  $h$  is admissible, then  $h(s_n) = h(s_n) = 0$ , and so  $f(p) = c(p)$ .
  - If  $p$  is sub-optimal, i.e.,  $c(p) \geq h^*(s_0)$ , then  $f(p) \geq h^*(s_0)$ .
  - Let  $p^* = \langle s_0^*, \dots, s_m^* \rangle$  where  $s_0^* = s_0$  and  $s_m^* = s_n$  be the optimal path to  $s_n$ .
  - A subpath of  $p^*$ , say  $p_{0:i}^*$  must still be on the frontier.
  - We compute  $f(p_{0:i}^*) = c(p_{0:i}^*) + h(s_i^*) \leq c(p_{0:i}^*) + h^*(s_i^*) = h^*(s_0) \leq f(p)$ .
  - However, this means that we should have explored  $p_{0:i}^*$  before  $p$ .


- **Time Complexity:**  $O(b^{c^*}/\epsilon)$ .
  - In the worst-case, the time-complexity of A\* is the same as UCS.
- **Space Complexity:**  $O(b^{c^*}/\epsilon)$ .
  - In the worst-case, the space-complexity of A\* is the same as UCS.

- Like IDDFS, we can implement an iterative version of A\*.
- We start with an  $f$ -limit of  $f_{\max} = h(s_0)$ , and repeatedly perform A\*, increasing the  $f$ -limit each time, until a solution is found
- In each iteration of IDA\*, we increase  $f_{\max}$  to the minimum  $f$ -value of all pruned paths in the previous iteration.
- This is called **iterative-deepening A-star search** (IDA\*).

# Iterative-Deepening A-Star Search: Example

- IDA\* is like performing A\* multiple times on a sub-tree of possible paths.


ltr.	$p$	$\mathcal{O}$
1	–	$A$
2	$A$	



# Iterative-Deepening A-Star Search: Example

- IDA\* is like performing A\* multiple times on a sub-tree of possible paths.

ltr.	$p$	$\mathcal{O}$
1	–	$A$
2	$A$	$AB$
3	$AB$	$ABC$
4	$ABC$	






# Iterative-Deepening A-Star Search: Example

- IDA\* is like performing A\* multiple times on a sub-tree of possible paths.

ltr.	$p$	$\mathcal{O}$
1	–	$A$
2	$A$	$AB, AC$
3	$AB$	$AC, ABC$
4	$ABC$	$AC$
5	$AC$	



- Below, we summarize the properties of GBFS, A\*, and IDA\*:

Property	GBFS	A*
Complete	$b, d < \infty$ , PC	$b < \infty, \varepsilon > 0$
Optimal	never <sup>2</sup>	$h$ admissible
Time Complexity	$O(b^m)$	$O(b^{c^*/\varepsilon})$
Space Complexity	$O(bm)$	$O(b^{c^*/\varepsilon})$

---

<sup>2</sup>except by chance.

- Weighted A\* is a generalization of the A\* algorithm where we let

$$f(\cdot) = wc(\cdot) + (1 - w)h(\cdot),$$

for some  $0 \leq w \leq 1$ .

- This way, we can interpret UCS as weighted A\* with  $w = 1$ , GBFS as weighted A\* with  $w = 0$ , and A\* as weighted A\* with  $w = 0.5$ .
- The weight,  $w$ , can be tuned to balance the trade-off between the time required to find a solution versus the quality of the solution found; increasing  $w$  results in a better solution, but smaller  $w$  results in a quicker solution.

- Designing heuristics is an art.
- There are many techniques used in practice. We shall consider two of them:
  - ① using a relaxation of the problem
  - ② using pattern databases

# Designing Heuristics via Problem Relaxation

- One way to design a heuristic for a problem is to find the solution of a relaxed problem and use its' cost as a heuristic for the original problem.

## Example: Slider Puzzle Relaxation ( $l_1$ distances)

- In the slider puzzle, tiles could only be moved into an unoccupied, horizontally/vertically adjacent space.
- Suppose we can move tiles to any horizontally/vertically adjacent space (including occupied ones).
- In this case, the sum of the  $l_1$  distances between each tile's correct position and its current position gives the number of moves required.
- For the example on the right, this value is  $3 + 2 + 0 + 1 + 1 + 2 + 0 + 1 = 10$ .

4	6	3
2		1
7	5	8

$s_0$

1	2	3
4	5	6
7	8	

$s \in \mathcal{G}$

# Designing Heuristics via Problem Relaxation

## Example: Slider Puzzle Relaxation (misplaced tiles)

- In the slider puzzle, tiles could only be moved into an unoccupied, horizontally/vertically adjacent space.
- Suppose we can move tiles to any adjacent space (including diagonally adjacent or occupied ones).
- In this case, the number of misplaced tiles gives the number of moves required.
- For the example on the right, this value is  $1 + 1 + 0 + 1 + 1 + 1 + 0 + 1 = 6$ .

4	6	3
2		1
7	5	8


$s_0$

1	2	3
4	5	6
7	8	

$s \in \mathcal{G}$

# Designing Heuristics via Pattern Databases

- Another way to design a heuristic for a problem is to break it up into sub-problem such that solving the original problem is equivalent to solving each of the sub-problems.
- **Example:** Slider Puzzle Pattern Database
  - It turns out that the Slider puzzle can be solved by first moving the '1' tile to the top-left position and then never moving it again.



For this particular example, it takes at least 11 moves to get from  $s_0$  to  $s_i$ .

# Designing Heuristics via Pattern Databases


- Suppose the costs are reversible, i.e., the cost to get from  $s_i$  to  $s_j$  is the same as the cost to get from  $s_j$  to  $s_i$ .
- For each  $s$ , we can search backwards from the solution(s) of the sub-problem until we find  $s$ . The cost of the resulting path is the cost of solving the sub-problem when the initial state is  $s$ .
- We store these values in a table. We can then compute a heuristic for the original problem using this table.



# Designing Heuristics via Pattern Databases (continued)

## Example: Slider Puzzle Pattern Database

- Suppose we had a function,  $h_{0,i}(s)$ , which gives the cost of getting from any state,  $s$ , to one of the form of  $s_i$ , and  $h_i$  is a heuristic for the sub-problem of finding  $g$  from  $s_i$ .



- We can then define  $h(p) = h_{0,i}(p) + h_i(p)$  as a heuristic for the original problem.

- In general, we can create many admissible heuristics for a given problem.
- However, some heuristics are better than others, in the sense that they will allow us to explore less paths.
- Given two admissible heuristics,  $h_1$ , and  $h_2$ , if  $h_1(s) \geq h_2(s)$  for all  $s$ , then we say that  $h_1$  **weakly dominates**  $h_2$ .
- If  $h_1(s) > h_2(s)$  for some  $s$ , then we say that  $h_1$  **strongly dominates**  $h_2$ .
- A strongly dominant heuristic is guaranteed to explore fewer paths.
- Clearly, if  $h_1$  and  $h_2$  are admissible heuristics, then  $h = \max\{h_1, h_2\}$  is also admissible and weakly dominates  $h_1$  and  $h_2$ .